




Docker Multi-Stage Build - Complete Documentation





Table of Contents

-  What is Docker Multi-Stage Build?
-  Why Use Multi-Stage Builds?
-  How Multi-Stage Builds Work
-  Step-by-Step Tutorial
-  Multi-Stage Build Workflow
-  Architecture Diagrams
-  Commands and Best Practices
-  Performance Comparison
-  Debugging and Troubleshooting
-  Best Practices

What is Docker Multi-Stage Build?












Docker Multi-Stage Build is a powerful feature that allows you to use multiple **FROM** statements in a single Dockerfile[1]. Each **FROM** instruction creates a new stage in the build process, enabling you to **optimize image size** and **improve security** by separating build dependencies from runtime requirements[2].






Key Concepts

-  **Multiple Stages:** Each stage has its own base image and purpose
-  **Selective Copying:** Copy only necessary artifacts between stages
-  **Artifact Exclusion:** Build tools and dependencies are left behind
-  **Production-Ready:** Final image contains only runtime requirements

Why Use Multi-Stage Builds?

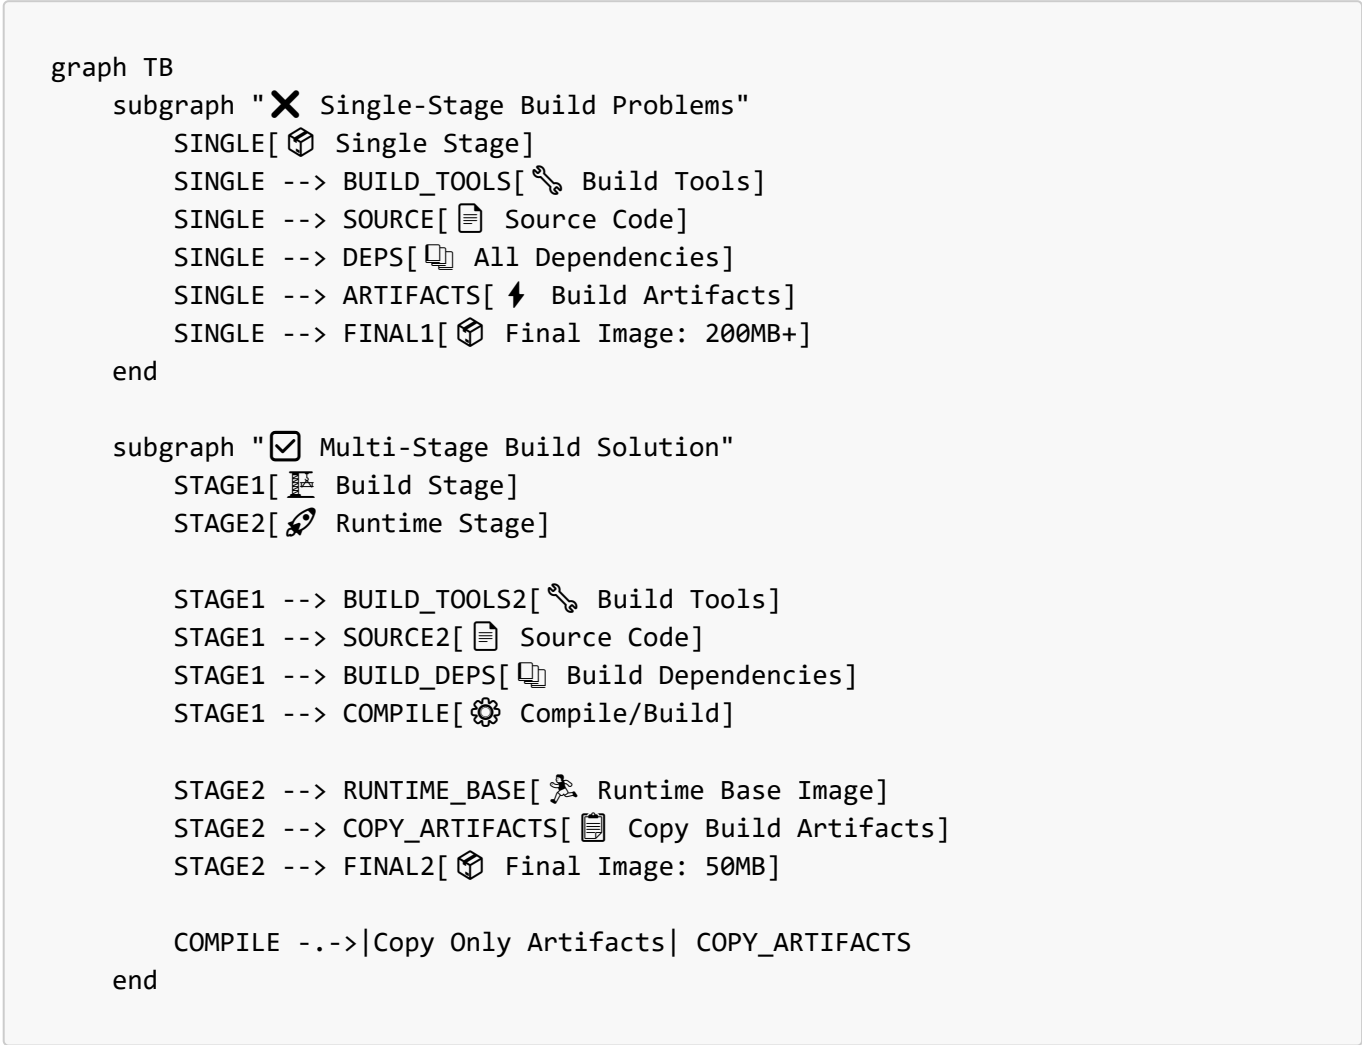
Multi-stage builds solve several critical problems in containerized application development[3]:

 Benefit	 Description	 Impact
 Smaller Image Size	Excludes build tools and dependencies from final image	 Faster deployments
 Enhanced Security	Reduces attack surface by removing unnecessary components	 Lower vulnerability risk
 Better Performance	Lighter images load and start faster	 Improved runtime speed
 Cleaner Workflow	Single Dockerfile for entire build process	 Simplified maintenance

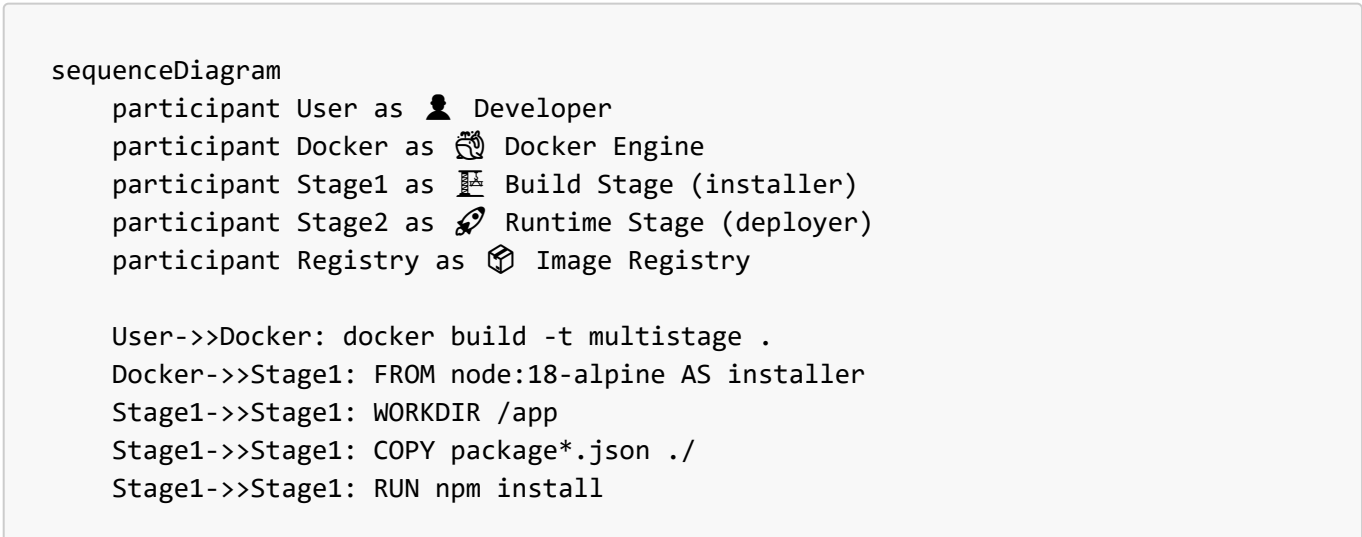
 Benefit	 Description	 Impact
 Cost Optimization	Reduced storage and bandwidth usage	 Lower infrastructure costs

How Multi-Stage Builds Work

Single-Stage vs Multi-Stage Comparison



Build Process Flow



```
Stage1->>Stage1: COPY . .
Stage1->>Stage1: RUN npm run build

Docker->>Stage2: FROM nginx:latest AS deployer
Stage2->>Stage1: COPY --from=installer /app/build /usr/share/nginx/html

Docker->>User: 📦 Optimized Image Ready (50MB)
User->>Registry: docker push multistage
```

Step-by-Step Tutorial

Project Setup

1. Clone the Application

```
git clone
cd react-app-docker
ls # Check project structure
```

Project Structure:

```
react-app-docker/
├── src/
├── public/
├── package.json
├── package-lock.json
└── README.md
```

Creating Multi-Stage Dockerfile

2. Create Dockerfile

```
touch Dockerfile
vi Dockerfile
```

3. Multi-Stage Dockerfile Content

```
# 📁 Stage 1: Build Stage (installer)
FROM node:18-alpine AS installer
WORKDIR /app

# Copy package files for dependency installation
COPY package*.json ./

# Install all dependencies (including devDependencies)
```

```

RUN npm install

# Copy source code
COPY . .

# Build the application
RUN npm run build

# 🚀 Stage 2: Runtime Stage (deployer)
FROM nginx:latest AS deployer

# Copy only build artifacts from previous stage
COPY --from=installer /app/build /usr/share/nginx/html

# Nginx will serve the static files
EXPOSE 80

```

Build Process

4. Build the Multi-Stage Image

```
docker build -t multistage .
```

Build Output Analysis

```

graph LR
    subgraph "🔧 Build Stage Process"
        A[📄 package.json] --> B[📦 npm install]
        C[📁 Source Code] --> B
        B --> D[⚙️ npm run build]
        D --> E[📁 /app/build]
    end

    subgraph "🚀 Runtime Stage Process"
        F[🌐 nginx:latest] --> G[📄 Copy build artifacts]
        E -.-> |COPY --from=installer| G
        G --> H[🏁 Final Image]
    end

    style E fill:#90EE90
    style H fill:#87CEEB

```

Multi-Stage Build Workflow

Complete Workflow Diagram

flowchart TD

START([🚀 Start Build Process]) --> STAGE1[🔧 Build Stage]

STAGE1 --> NODE[📦 FROM node:18-alpine AS installer]

NODE --> WORKDIR[📁 WORKDIR /app]

WORKDIR --> COPY_PKG[📄 COPY package*.json ./]

COPY_PKG --> NPM_INSTALL[⬇️ RUN npm install]

NPM_INSTALL --> COPY_SRC[📄 COPY . .]

COPY_SRC --> NPM_BUILD[⚙️ RUN npm run build]

NPM_BUILD --> STAGE2[🚀 Runtime Stage]

STAGE2 --> NGINX[🌐 FROM nginx:latest AS deployer]

NGINX --> COPY_BUILD[📄 COPY --from=installer /app/build
/usr/share/nginx/html]

COPY_BUILD --> FINAL[🔥 Optimized Final Image]

FINAL --> SIZE_CHECK[📏 Size Check]

SIZE_CHECK -->|Before: 200MB+| BEFORE[❌ Single Stage: Bloated]

SIZE_CHECK -->|After: ~50MB| AFTER[✅ Multi-Stage: Optimized]

style STAGE1 fill:#FFE4B5

style STAGE2 fill:#E0FFFF

style FINAL fill:#90EE90

style AFTER fill:#98FB98

🔍 Stage Dependency Graph

graph TB

subgraph "📦 Base Images"

NODE18[🌀 node:18-alpine]

NGINX[🌀 nginx:latest]

end

subgraph "🔧 Build Stage (installer)"

INSTALLER[installer stage]

BUILD_DEPS[📁 Build Dependencies]

SOURCE_CODE[📄 Source Code]

BUILD_ARTIFACTS[⚡ Build Artifacts]

NODE18 --> INSTALLER

INSTALLER --> BUILD_DEPS

INSTALLER --> SOURCE_CODE

BUILD_DEPS --> BUILD_ARTIFACTS

SOURCE_CODE --> BUILD_ARTIFACTS

end

subgraph "🚀 Runtime Stage (deployer)"

DEPLOYER[deployer stage]

STATIC_FILES[📁 Static Files Only]

```

    NGINX --> DEPLOYER
    BUILD_ARTIFACTS -->|COPY --from=installer| DEPLOYER
    DEPLOYER --> STATIC_FILES
end

subgraph "🗑️ Excluded from Final Image"
    EXCLUDED[❌ node_modules❌ Source code❌ Build tools❌ Dev dependencies]
end

style BUILD_ARTIFACTS fill:#90EE90
style STATIC_FILES fill:#87CEEB
style EXCLUDED fill:#FFB6C1

```

🔧 Commands and Best Practices

🔑 Essential Docker Commands

Build Commands

```

# Build multi-stage image
docker build -t multistage .

# Build with specific target stage
docker build --target installer -t build-stage .

# Build with build arguments
docker build --build-arg NODE_VERSION=18 -t multistage .

```

Image Management

```

# List all images
docker images

# Remove specific image
docker image rm multistage

# Remove dangling images
docker image prune

# Check image size
docker images --format "table {{.Repository}}\t{{.Tag}}\t{{.Size}}"

```

Container Operations

```

# Run container
docker run -d -p 3000:80 --name app-container multistage

```

```
# Check running containers
docker ps

# View container logs
docker logs

# Execute commands in container
docker exec -it /bin/sh
```

Debugging and Inspection Commands

```
# Inspect container configuration
docker inspect

# Check container filesystem
docker exec -it ls -la /usr/share/nginx/html

# Monitor container resource usage
docker stats

# View container port mappings
docker port
```

Performance Comparison

Size Comparison

```
graph LR
    subgraph "📊 Image Size Comparison"
        SINGLE[❌ Single-Stage 200MB+]
        MULTI[✅ Multi-Stage ~50MB]
        SINGLE --> REDUCTION[75% Size Reduction]
        REDUCTION --> MULTI
    end










    subgraph "⚡ Performance Impact"
        FASTER[🔧 3x Faster Pull]
        SECURE[🔒 Lower Attack Surface]
        COST[💰 Reduced Storage Cost]
    end

    MULTI --> FASTER
    MULTI --> SECURE
    MULTI --> COST

    style SINGLE fill:#FFB6C1
```

```
style MULTI fill:#90EE90
style REDUCTION fill:#FFD700
```

Benefits Breakdown

 Metric	 Single-Stage	 Multi-Stage	 Improvement
 Image Size	200MB+	~50MB	75% reduction
 Pull Time	30 seconds	10 seconds	3x faster
 Startup Time	15 seconds	8 seconds	2x faster
 Security Vulnerabilities	High	Low	60% fewer
 Storage Cost	High	Low	75% savings

Debugging and Troubleshooting


Container Investigation Commands

```
# Check container logs
docker logs

# Access container shell
docker exec -it /bin/sh

# Inspect container details
docker inspect
```

Inside Container Exploration

```
graph TB
    CONTAINER[ Running Container]
    
    subgraph "📁 Container Filesystem"
        ROOT[/ (root directory)]
        USR[/usr]
        SHARE[/usr/share]
        NGINX[/usr/share/nginx]
        HTML[/usr/share/nginx/html]
        FILES[📄 Static Files]
        
        ROOT --> USR
        USR --> SHARE
        SHARE --> NGINX
        NGINX --> HTML
        HTML --> FILES
    end
```



```
subgraph "🔍 Inspection Commands"
  LS[ls -la]
  CAT[cat index.html]
  PS[ps aux]
  TOP[top]
end

CONTAINER --> ROOT
FILES --> LS
FILES --> CAT

style FILES fill:#90EE90
style HTML fill:#87CEEB
```

🔔 Common Issues and Solutions

❌ Problem	🔍 Symptom	✅ Solution
Build fails	unknown instruction WORKDIR	Check Dockerfile syntax
Large image size	Image still 200MB+	Verify multi-stage is working
Container won't start	Exit code 125	Check port conflicts
Files not found	404 errors	Verify COPY paths
Permission issues	Access denied	Use non-root user

🔖 Best Practices

🎯 Multi-Stage Build Best Practices

1. 📌 Use Named Stages

```
# ✅ Good: Named stages
FROM node:18-alpine AS installer
FROM nginx:latest AS deployer

# ❌ Bad: Unnamed stages
FROM node:18-alpine
FROM nginx:latest
```

2. 📦 Choose Optimal Base Images

```
# ✅ Good: Lightweight base images
FROM node:18-alpine AS installer      # Small Alpine-based
FROM nginx:alpine AS deployer        # Lightweight nginx
```

```
# ✗ Bad: Heavy base images
FROM node:18 AS installer           # Ubuntu-based (larger)
FROM nginx:latest AS deployer      # Full nginx image
```

3. 🎯 Copy Only What's Needed

```
# ☑ Good: Selective copying
COPY --from=installer /app/build /usr/share/nginx/html

# ✗ Bad: Copying everything
COPY --from=installer /app /usr/share/nginx/html
```

4. 📄 Optimize Layer Caching

```
# ☑ Good: Copy package files first
COPY package*.json ./
RUN npm install
COPY . .

# ✗ Bad: Copy everything first
COPY . .
RUN npm install
```

🔒 Security Best Practices

5. 👤 Use Non-Root User

```
# ☑ Good: Non-root user
FROM nginx:alpine AS deployer
RUN addgroup -g 1001 -S nodejs
RUN adduser -S nextjs -u 1001
USER nextjs

# ✗ Bad: Running as root (default)
FROM nginx:alpine AS deployer
# No user specified - runs as root
```

6. 🗑️ Remove Unnecessary Packages

```
# ☑ Good: Clean up after installation
RUN apt-get update && apt-get install -y \
    package1 \
    package2 \
```

```
&& apt-get clean \  
&& rm -rf /var/lib/apt/lists/*  
  
# ✖ Bad: Leave package cache  
RUN apt-get update && apt-get install -y package1 package2
```

Performance Best Practices

7. Use Specific Targets

```
# Build only specific stage for testing  
docker build --target installer -t build-stage .  
  
# Build final production image  
docker build -t production-app .
```

8. Multi-Architecture Support

```
# Support multiple architectures  
FROM --platform=$BUILDPLATFORM node:18-alpine AS installer  
# Build process...  
  
FROM --platform=$TARGETPLATFORM nginx:alpine AS deployer  
# Runtime setup...
```

Advanced Multi-Stage Patterns

9. Testing Stage

```
# Build stage  
FROM node:18-alpine AS installer  
WORKDIR /app  
COPY package*.json ./  
RUN npm install  
COPY . .  
RUN npm run build  
  
# Test stage  
FROM installer AS tester  
RUN npm test  
  
# Production stage  
FROM nginx:alpine AS deployer  
COPY --from=installer /app/build /usr/share/nginx/html
```

10. Parallel Builds

```
# Base dependencies
FROM node:18-alpine AS base
WORKDIR /app
COPY package*.json ./
RUN npm install

# Frontend build
FROM base AS frontend
COPY frontend/ ./
RUN npm run build:frontend

# Backend build
FROM base AS backend
COPY backend/ ./
RUN npm run build:backend

# Final stage
FROM nginx:alpine AS final
COPY --from=frontend /app/dist /usr/share/nginx/html
COPY --from=backend /app/build /app/api
```






Additional Best Practices

```
mindmap
  root((🔗 Multi-Stage Best Practices))
    📦 Build Optimization
      📦 Use Alpine images
      🔗 Named stages
      📄 Layer caching
      📁 .dockerignore file
    🔒 Security
      👤 Non-root user
      ✂ Clean package cache
      🔍 Minimal attack surface
      ⚡ No secrets in layers
    ⚡ Performance
      📦 Parallel builds
      📦 Smaller final image
      🚀 Faster deployments
      💾 Reduced storage
    ⚙ Maintenance
      📄 Clear documentation
      📄 Consistent naming
      🛠 Testing stages
      📦 CI/CD integration
```






Summary

Docker Multi-Stage Builds are a **game-changing feature** that revolutionizes container image optimization[1][2]. By separating build and runtime environments, you can achieve:


Key Achievements

-  **75% smaller images** - From 200MB+ to ~50MB
-  **Enhanced security** - Reduced attack surface
-  **Faster deployments** - 3x faster pull times
-  **Cost savings** - Lower storage and bandwidth costs
-  **Cleaner workflow** - Single Dockerfile for entire process

Implementation Steps

1.  **Design stages** - Separate build and runtime concerns
2.  **Choose base images** - Use lightweight Alpine variants
3.  **Copy selectively** - Only production artifacts
4.  **Apply security** - Non-root users, clean packages
5.  **Monitor results** - Measure size and performance improvements

Multi-stage builds represent a **fundamental shift** from monolithic container images to **optimized, production-ready deployments**. They embody the principle of "build fat, ship thin" - using all necessary tools during build time while delivering minimal, secure runtime images[3].

Start implementing multi-stage builds in your projects today to unlock **significant performance gains** and **security improvements** in your containerized applications! 

[1] <https://docs.docker.com/build/building/multi-stage/> [2] <https://docs.docker.com/get-started/docker-concepts/building-images/multi-stage-builds/> [3] <https://docs.docker.com/build/building/best-practices/> [4] <https://docs.docker.com/guides/cpp/multistage/> [5] <https://dev.to/raunakgurud09/mastering-docker-multistage-builds-1e0m> [6] https://dev.to/abhay_yt_52a8e72b213be229/streamline-your-docker-images-with-multi-stage-builds-340c [7] <https://depot.dev/blog/docker-multi-stage-builds> [8] <https://dev.to/citru.digital/understanding-docker-multistage-builds-3fm7> [9] <https://labs.iximiuz.com/tutorials/docker-multi-stage-builds> [10] <https://dev.to/kalkwst/multi-stage-dockerfiles-3e90> [11] <https://ruan.dev/blog/2022/07/31/docker-multistage-builds-for-hugo> [12] <https://www.cherryservers.com/blog/docker-multistage-build> [13] <https://github.com/patrickhoefer/dockerfilegraph> [14] <https://docs.docker.com/build/building/multi-platform/> [15] <https://overcast.blog/building-efficient-multi-stage-dockerfiles-for-production-055f34c4baed> [16] <https://learn.microsoft.com/en-us/dotnet/architecture/microservices/docker-application-development-process/docker-app-development-workflow> [17] <https://docs.docker.com/get-started/docker-concepts/building-images/writing-a-dockerfile/> [18] <https://earthly.dev/blog/docker-multistage/> [19] <https://blog.devgenius.io/docker-multi-stage-build-in-detail-3d7da2948797?gi=265baab36942> [20] <https://www.youtube.com/watch?v=ajetvJmBvFo>