



Kubernetes Complete Documentation



Table of Contents

- [What is Kubernetes?](#)
- [History and Evolution](#)
- [Kubernetes Architecture](#)
- [Control Plane Components](#)
- [Worker Node Components](#)
- [Kubernetes Workflow](#)
- [Architecture Diagrams](#)
- [Key Benefits](#)



What is Kubernetes?

Kubernetes (also known as **K8s**) is an open-source container orchestration system that automates the deployment, scaling, and management of containerized applications. The name comes from the Greek word for "helmsman" - the person who steers a ship.[1][2]



Why K8s?

- **K** (first letter) + **8** (eight letters in between) + **s** (last letter) = **K8s**



Core Purpose

Kubernetes solves the **container orchestration** problem by automating:


- **Deployment** of containers
- **Scaling** applications up and down
- **Management** of container lifecycle
- **Self-healing** capabilities



History and Evolution



Traditional Deployment Era




Challenges:

- **Expensive** hardware procurement
- **Manual** environment setup

- 📊 **Poor** scalability
- 🚫 **Vendor lock-in**


☁️ Cloud Revolution (AWS Era)



Benefits:

- ⚡ **Quick** resource provisioning
- 🔄 **Auto-scaling** capabilities
- 🛠️ **Managed services** (RDS, ELB, etc.)
- 💵 **Pay-as-you-use** model

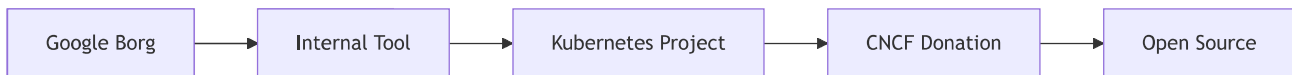
📦 Containerization Revolution



Evolution:

- 🏗️ **Heavy VMs** → 🍃 **Lightweight containers**
- 🚢 **Docker** made containerization accessible
- 🎯 **Need** for container orchestration emerged

🎯 Google's Solution: Borg → Kubernetes



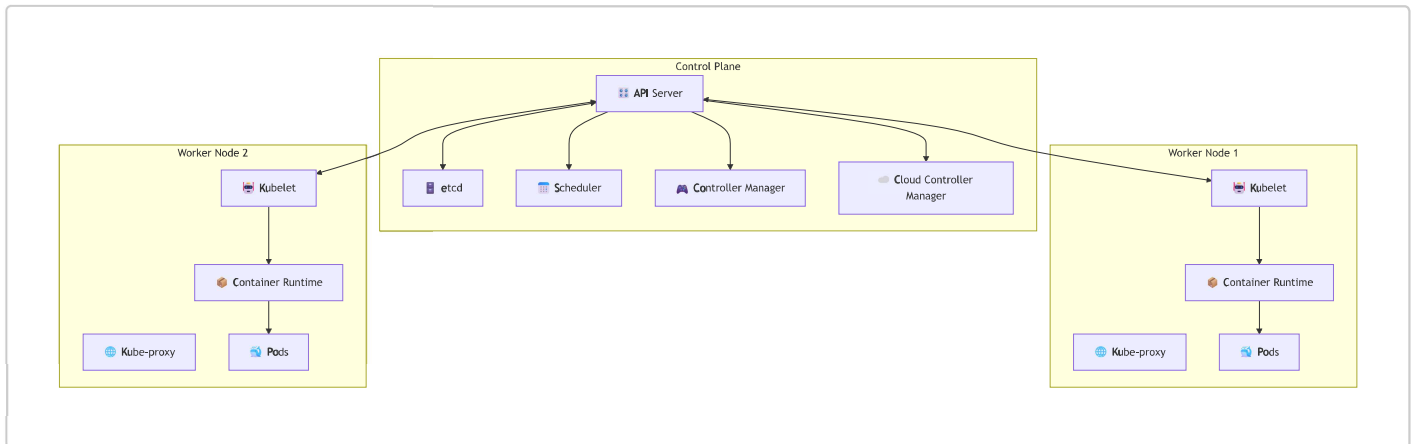
Timeline:

- 🏢 **Google** created **Borg** for internal use
- 🔄 **2014**: Kubernetes project started (ground-up rewrite)
- 📁 **2014**: Donated to **CNCF** (Cloud Native Computing Foundation)[2]

🏗️ Kubernetes Architecture

Kubernetes follows a **master-worker** architecture with two main components:[2]


High-Level Architecture







Control Plane Components

The **Control Plane** manages the overall state of the cluster. It consists of:[2][3]

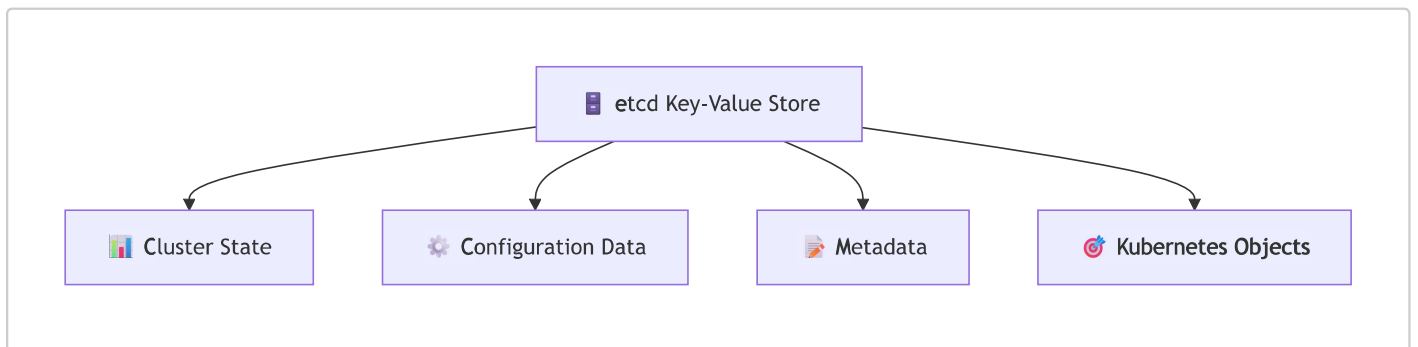
API Server



Functions:

-  **Entry point** for all administrative tasks
-  **Authentication** and **authorization**
-  **Validates** API requests
-  **Communication hub** between components[3]

etcd




Purpose:

-  **Distributed key-value store**

- 📅 Stores all **cluster state** information
- ⚙️ Contains **configuration data**
- 🔒 **Only accessible** via API Server[3]

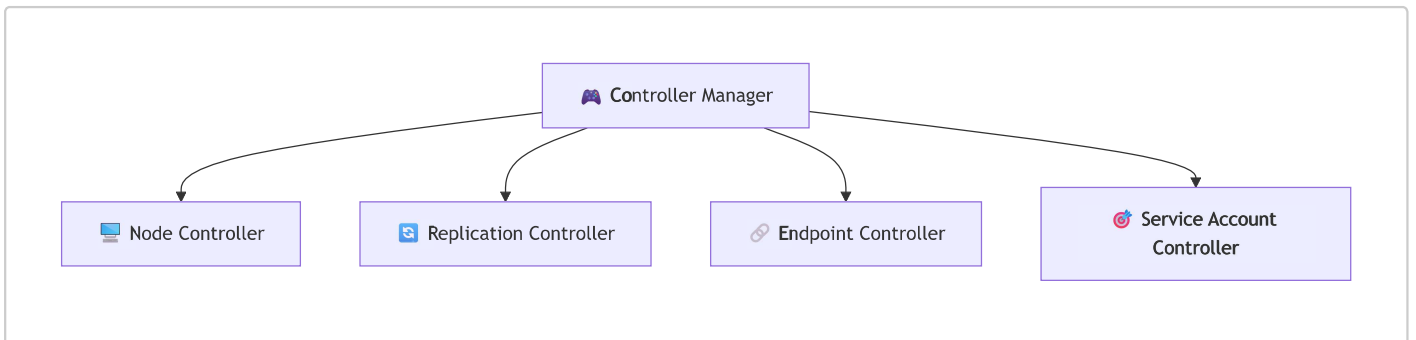
📅 Scheduler



Responsibilities:

- 📍 **Assigns pods** to appropriate worker nodes
- 🔍 **Evaluates** resource requirements
- ⚖️ **Load balancing** across nodes
- 📊 **Optimizes** resource utilization[2]

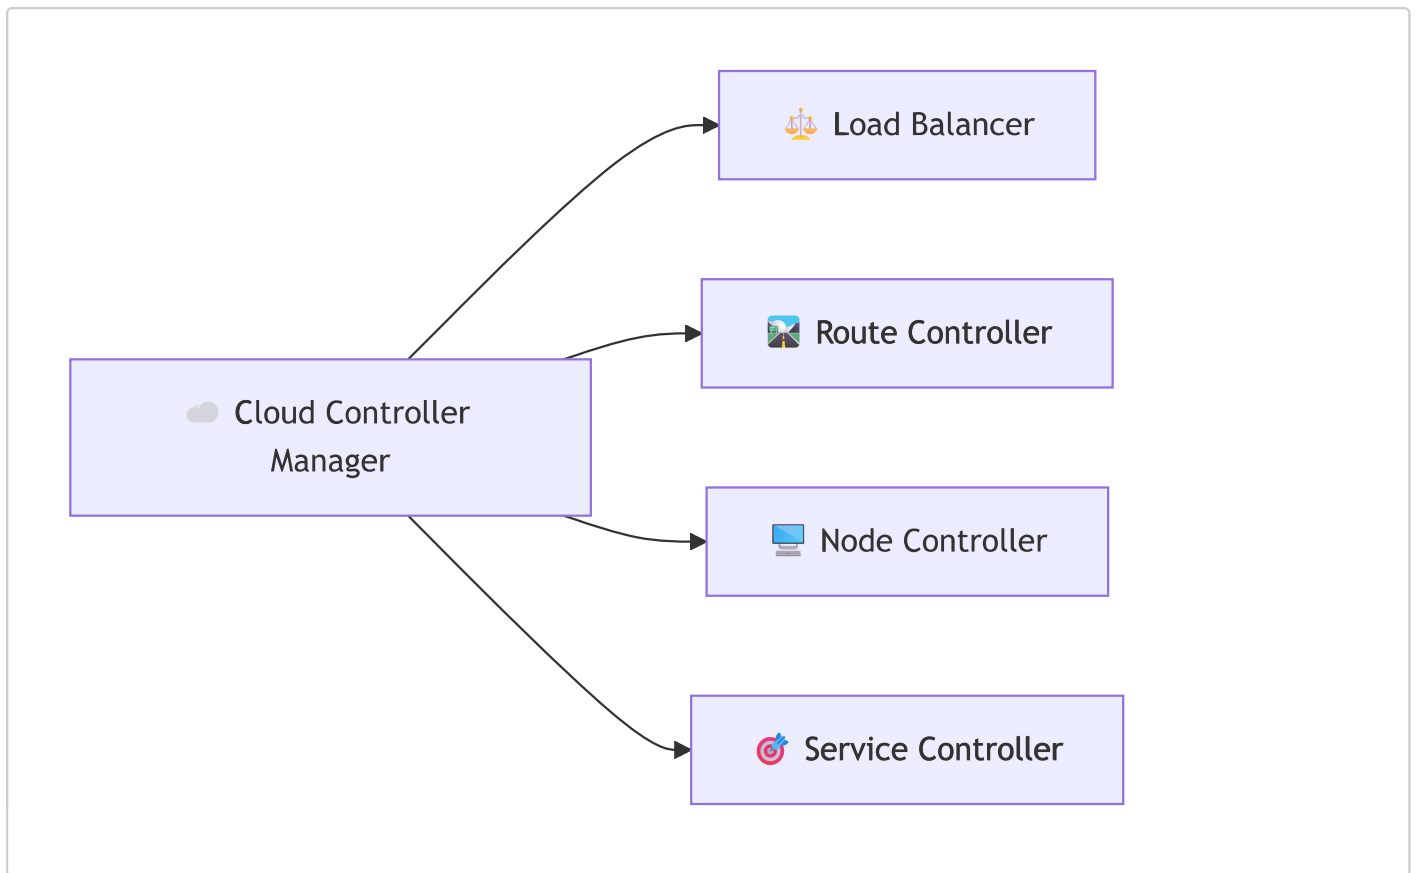
🎮 Controller Manager



Controllers Include:

- 💻 **Node Controller**: Monitors node health
- 🔄 **Replication Controller**: Manages pod replicas
- 🔗 **Endpoint Controller**: Updates endpoint objects
- 🎯 **Service Account Controller**: Creates default service accounts[2]

☁️ Cloud Controller Manager (CCM)

**Purpose:**

- ☁ **Cloud-specific** control logic
- ⚖ Manages **load balancers**
- 🗺 Sets up **network routes**
- 🔗 **Links cluster** to cloud provider APIs[2]

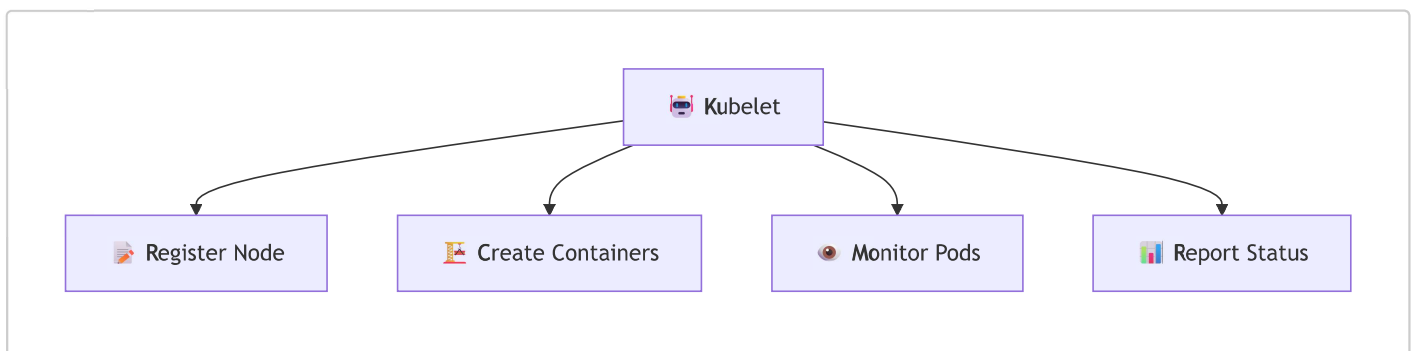


Worker Node Components

Worker Nodes run the actual containerized applications. Each node contains:[3]



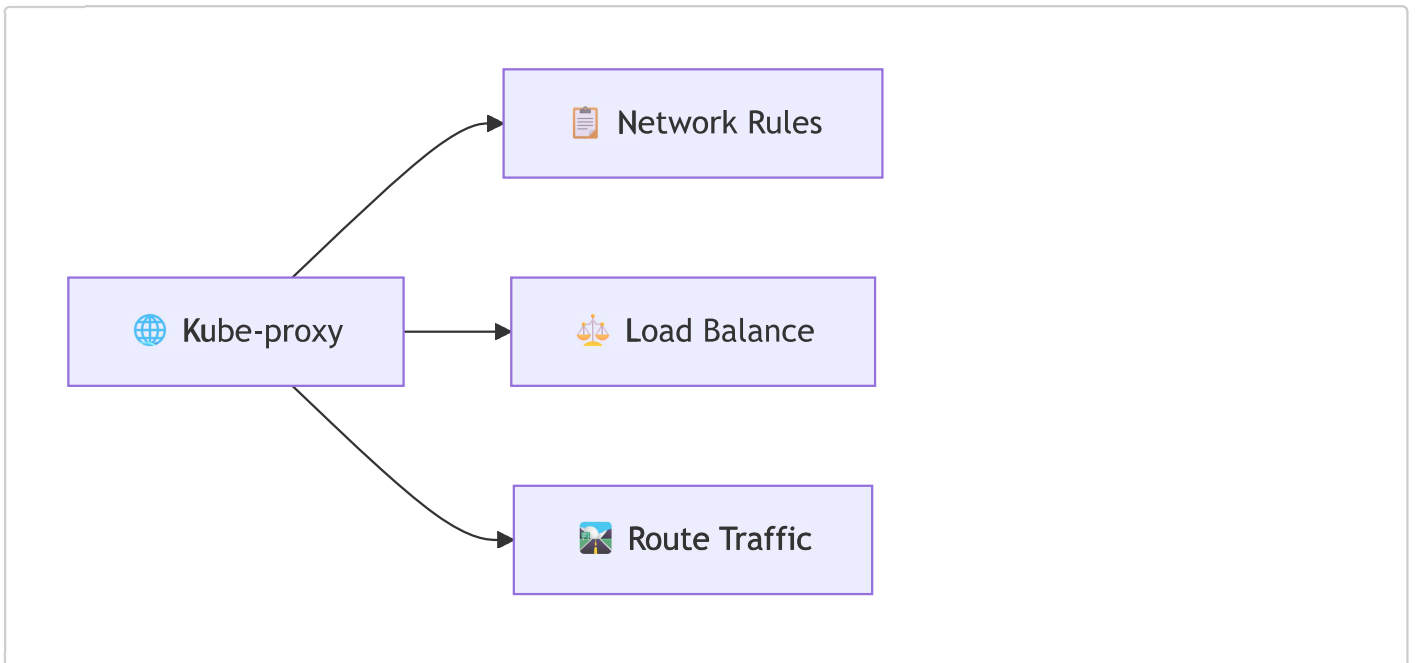
Kubelet

**Functions:**

- 📄 **Registers** worker node with API server
- 🏗 **Creates/manages** containers for pods

- 📺 **Monitors** pod health (liveness, readiness probes)
- 📊 **Reports** node and pod status[4]

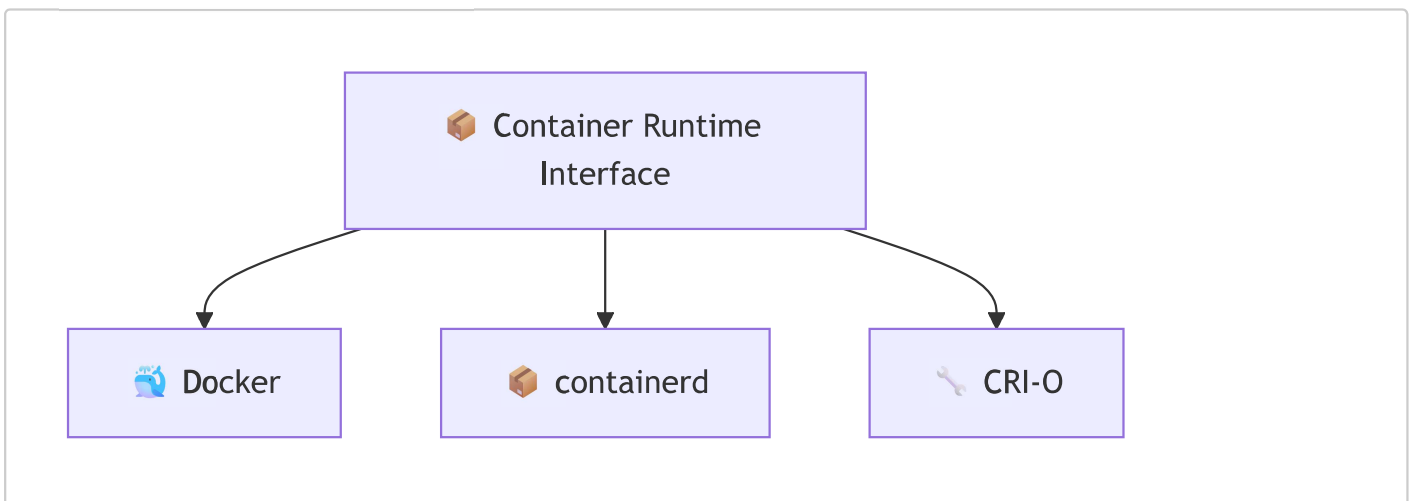
🌐 Kube-proxy



Responsibilities:

- 🌐 **Network proxy** on each node
- 📋 Maintains **network rules**
- ⚖️ **Load balances** traffic to pods
- 🗺️ **Routes** requests to appropriate pods[3]

📦 Container Runtime Interface (CRI)

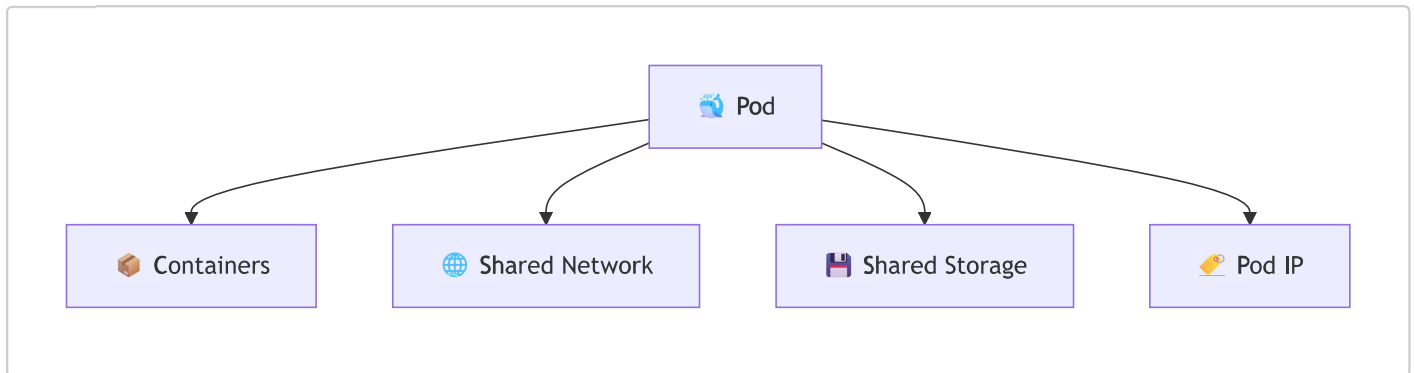


Options:

- 🐳 **Docker Engine**
- 📦 **containerd**

- 🛠️ **CRI-O**
- 🚀 **Other OCI-compliant runtimes**[4]

🐳 Pods

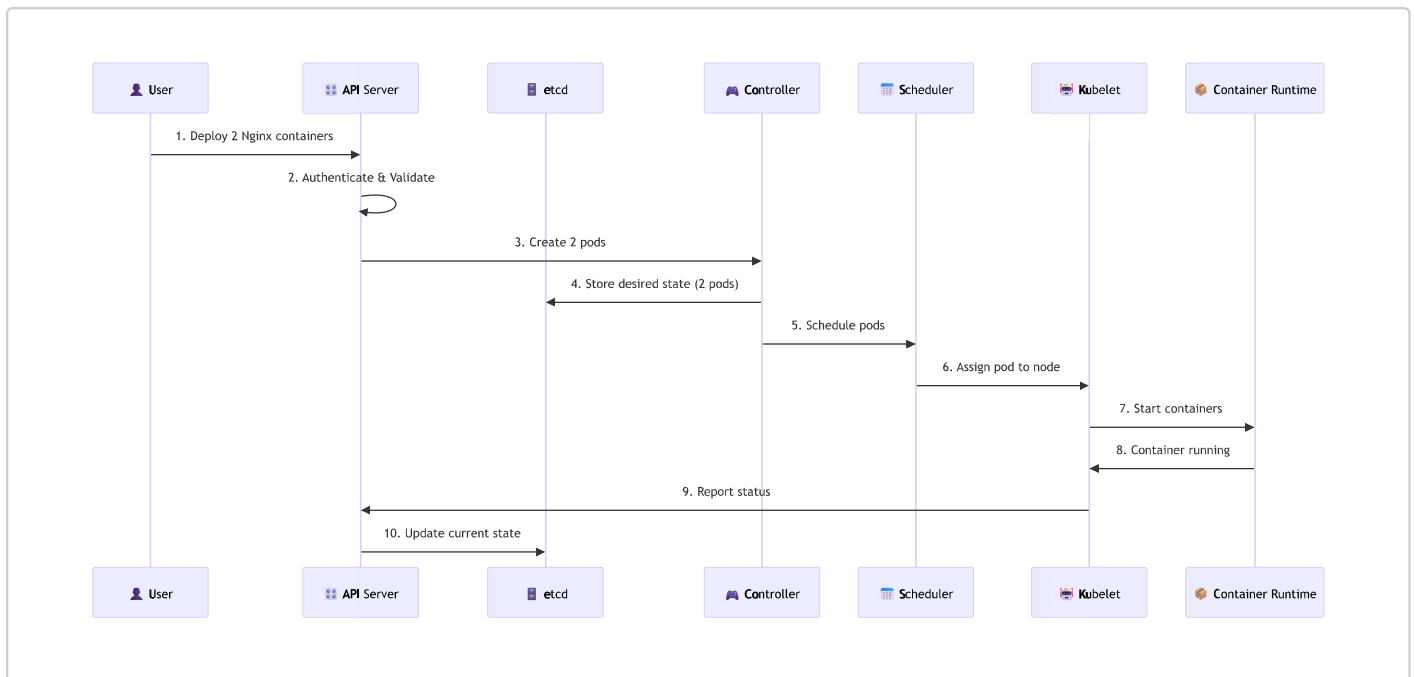


Characteristics:

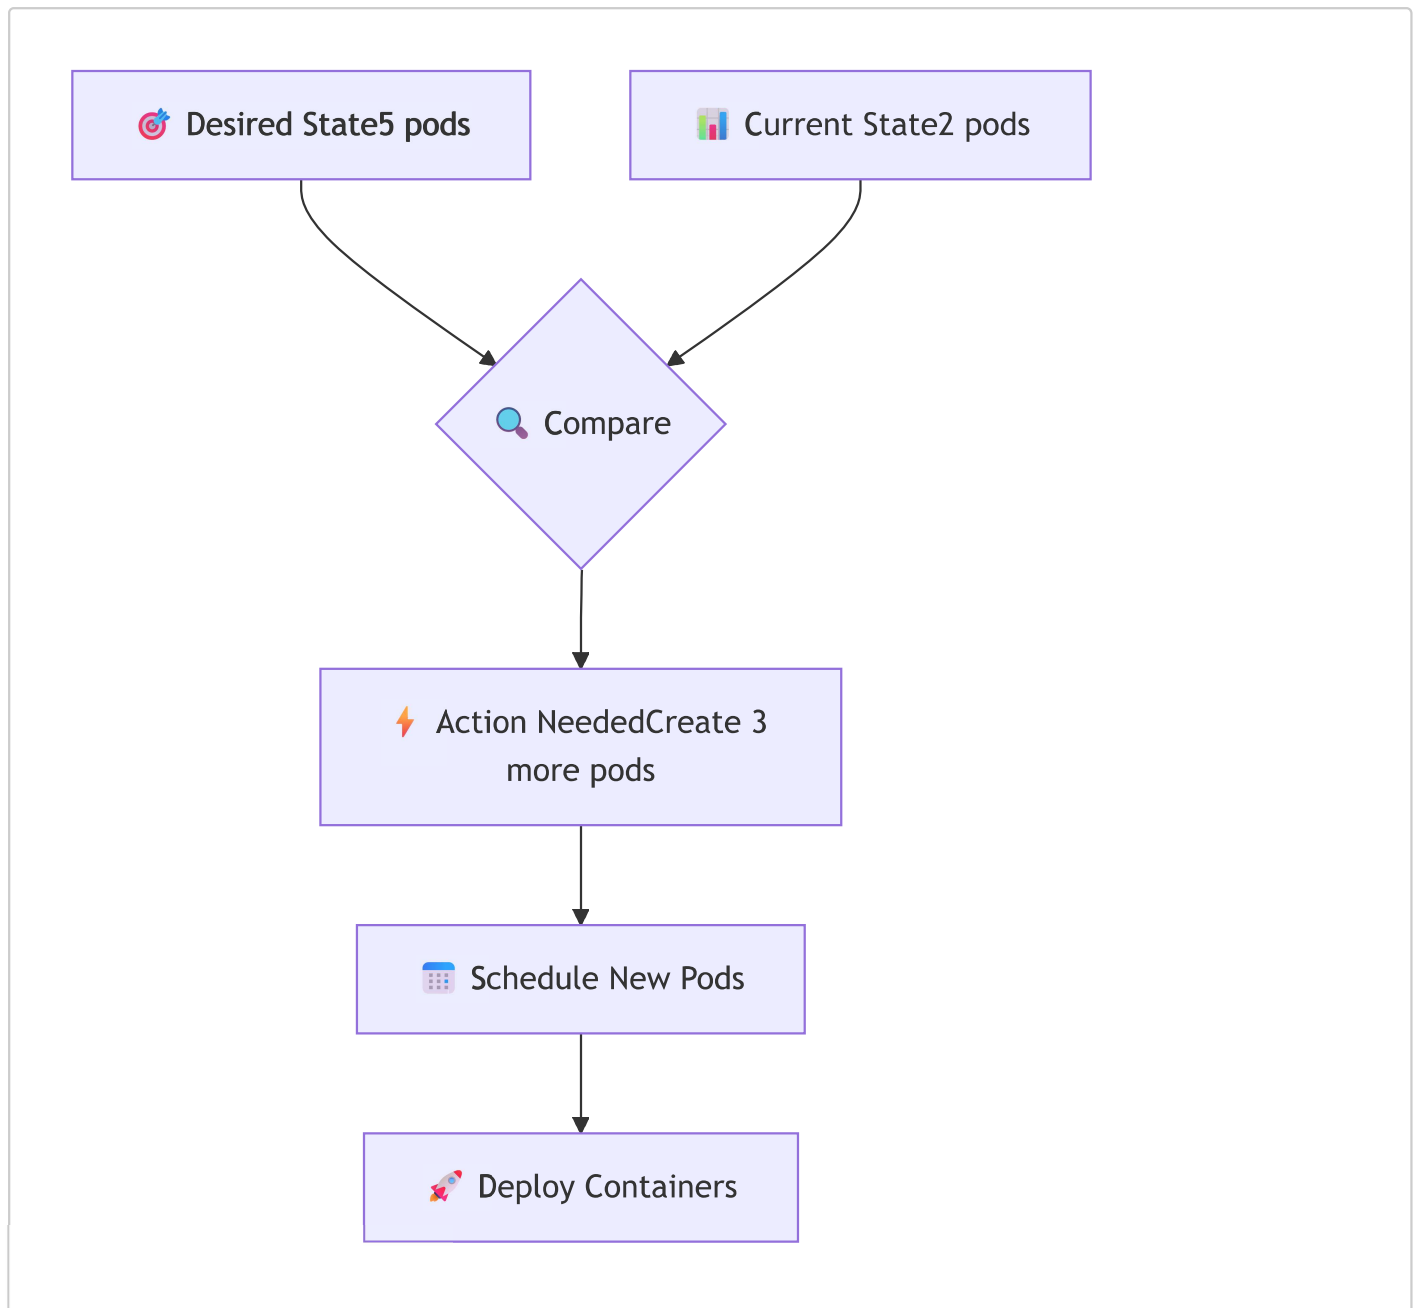
- 🐳 **Smallest deployable unit**
- 📦 Contains **one or more containers**
- 🌐 **Shared network** and storage
- 🏷️ Has **unique IP address**[3]

🔄 Kubernetes Workflow

📄 Deployment Process



🔄 State Reconciliation

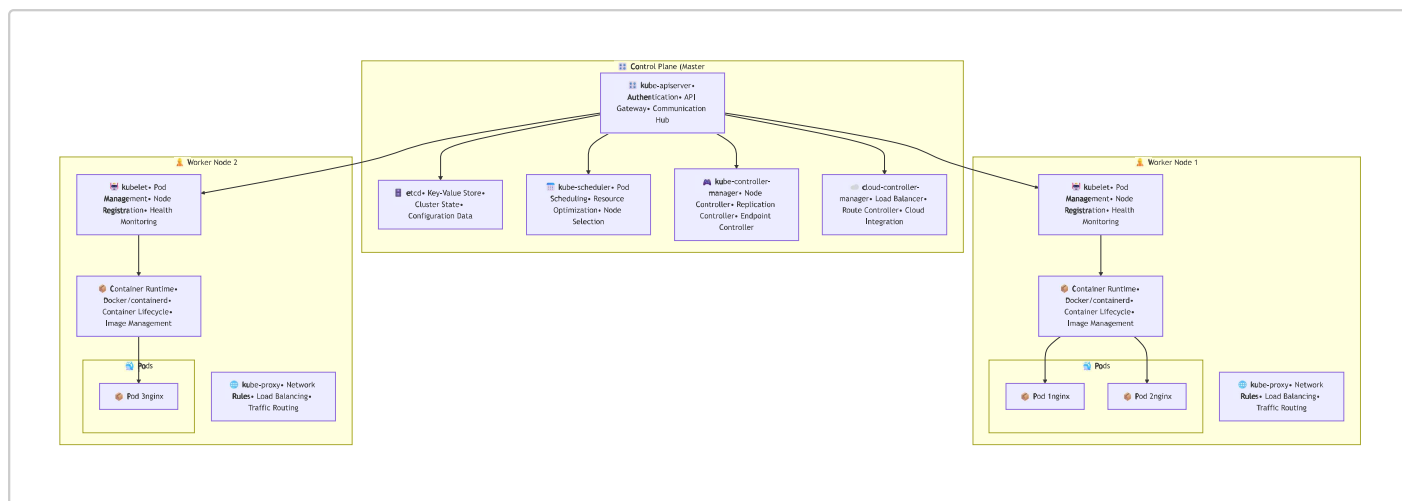


Key Concept: Declarative Management

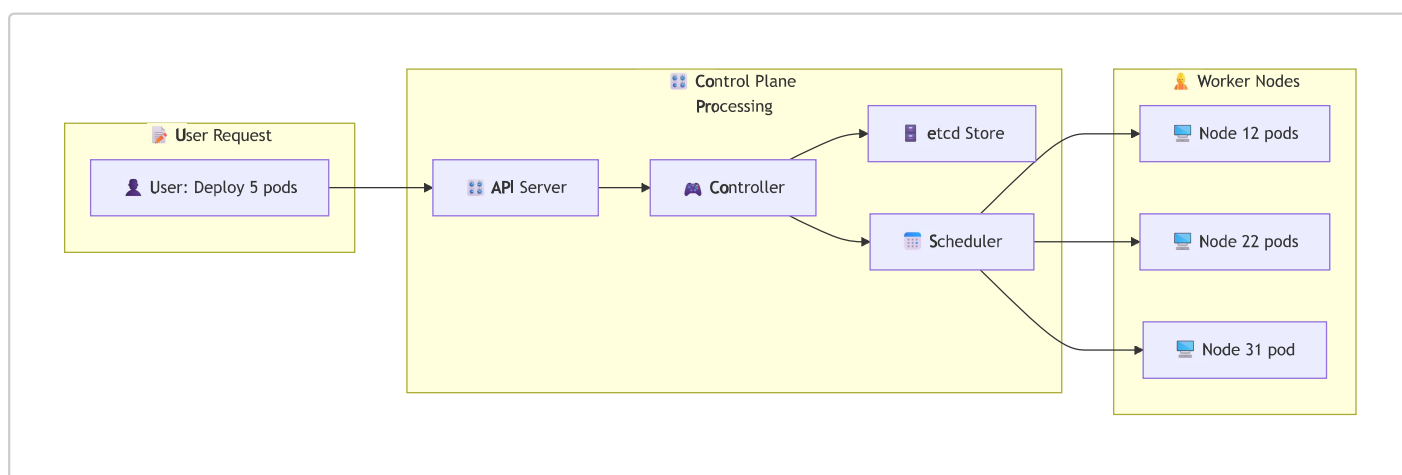
- 🎯 You define **desired state**
- 🔍 Kubernetes **continuously monitors**
- ⚡ **Automatically reconciles** differences
- 🔄 **Self-healing** capabilities

📊 Architecture Diagrams

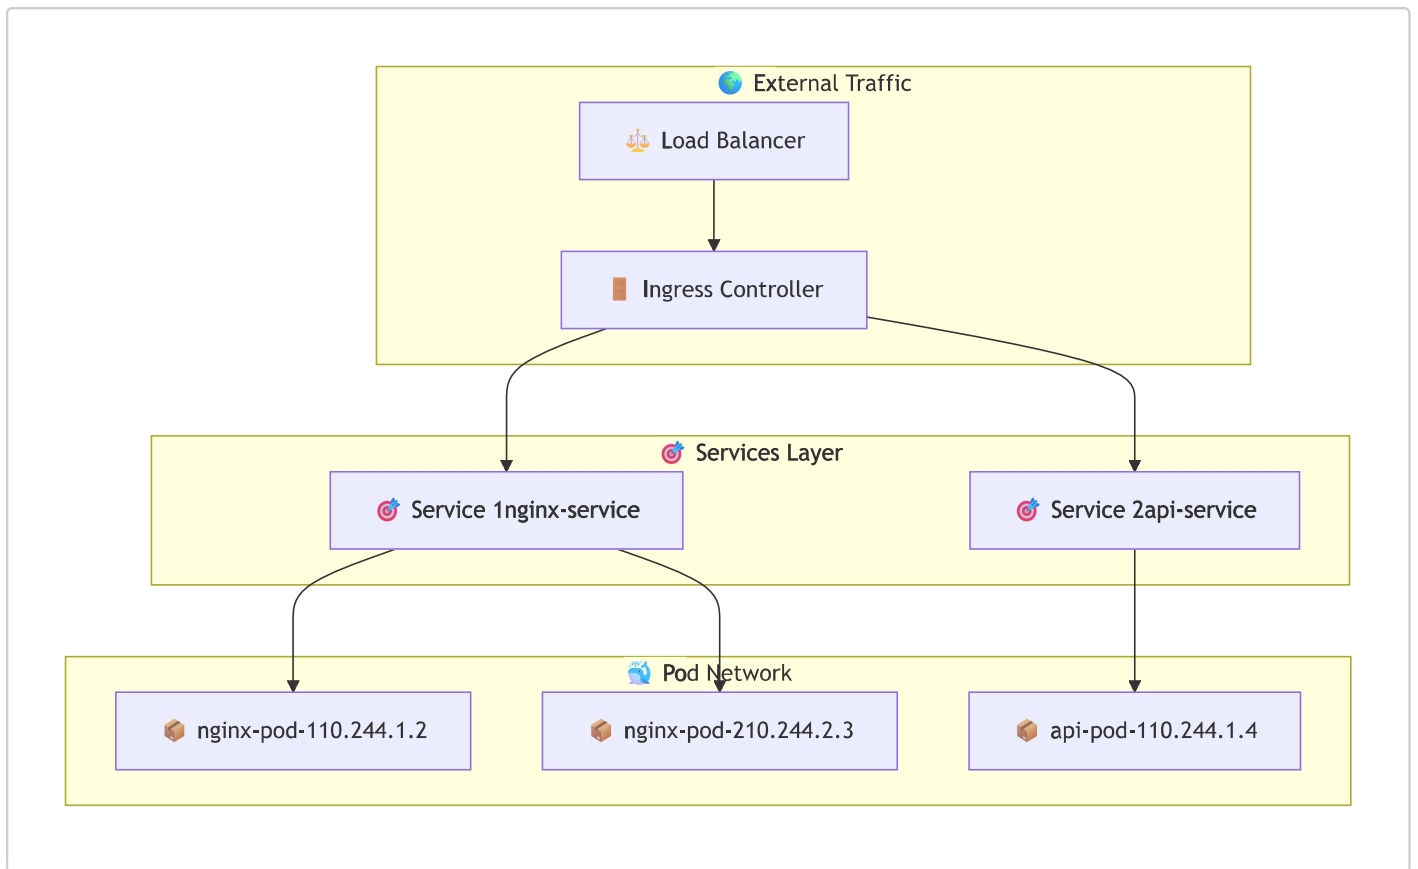
🏗️ Complete Kubernetes Architecture



Pod Lifecycle Flow



Networking Architecture

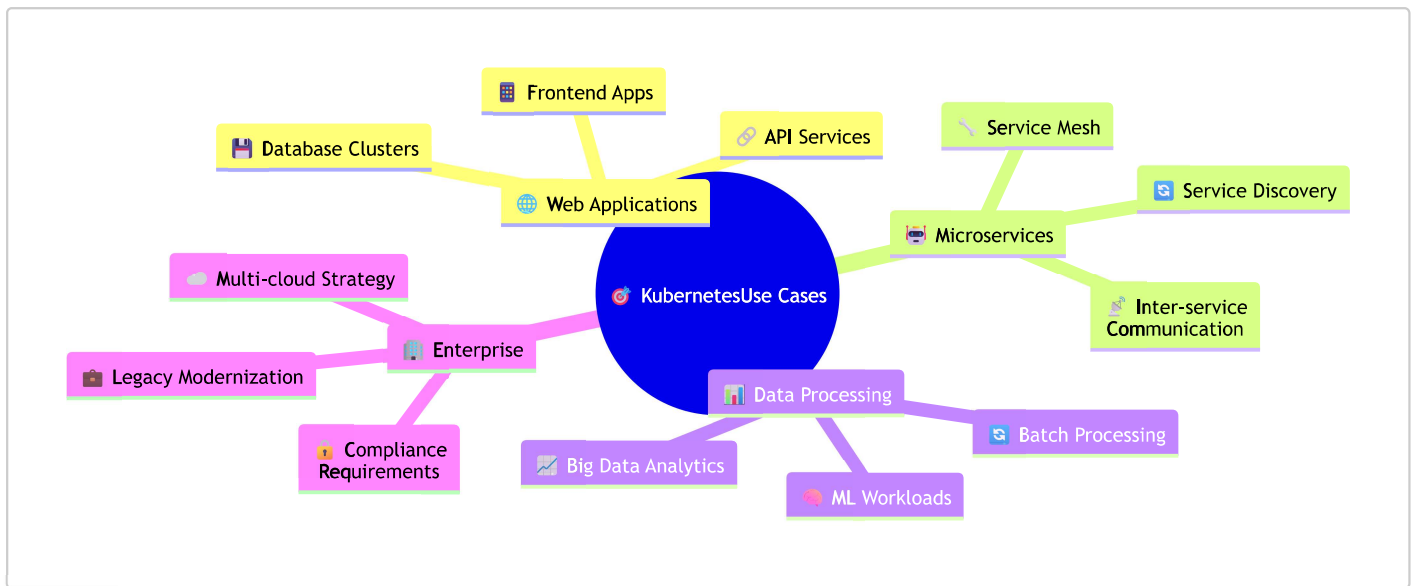


💡 Key Benefits


🚀 Why Choose Kubernetes?

Feature	Description	Benefit
🌐 Cloud Agnostic	Works on any infrastructure	🔒 No vendor lock-in
📈 Auto-scaling	Scales based on demand	💰 Cost optimization
🔄 Self-healing	Automatically replaces failed containers	🛡️ High availability
⚖️ Load Balancing	Distributes traffic efficiently	🚀 Better performance
🔒 Security	Built-in security features	🏢 Enterprise-ready
📦 Rolling Updates	Zero-downtime deployments	🔄 Continuous delivery

🎯 Use Cases



🔧 Ecosystem Integration



🎓 Summary

Kubernetes has revolutionized container orchestration by providing:

🔑 Key Takeaways

- **Centralized control** through the Control Plane
- **Distributed execution** via Worker Nodes
- **Declarative management** with desired state reconciliation
- **Cloud-agnostic** architecture preventing vendor lock-in
- **Enterprise-grade** features for production workloads

🚀 Getting Started

1. **Learn** container fundamentals (Docker)
2. **Understand** Kubernetes architecture
3. **Practice** with local clusters (minikube, kind)
4. **Deploy** to cloud providers (EKS, GKE, AKS)
5. **Master** kubectl and YAML manifests